BUSINESS PROCESS EVOLUTION
AT JPMORGAN CHASE & CO.

Analysis of Two Core Business Processes:

KYC / Customer Onboarding & Fraud Detection

Course: Business Process Engineering (BPE)

Submitted by: Waleed Luqman
Reg. No: FA22-BSE-119

Submitted to: Sir. Mukhtiar Zamin

COMSATS University Islamabad, Abbottabad Campus
25 May 2026

1. Introduction

JPMorgan Chase & Co. is one of the largest financial institutions in the world, with total assets exceeding $3.9 trillion and operations spanning more than 100 countries. As a global banking powerhouse, JPMorgan Chase depends on the continuous improvement of its core business processes to remain competitive, compliant, and customer-centric in an ever-evolving financial landscape.

This report presents a versioned analysis of two of the bank's most critical business processes: the KYC and Customer Onboarding process, and the Fraud Detection process. Rather than simply contrasting the original and current states, this report traces the complete evolutionary journey of each process across three clearly defined versions V1 representing the original manual or rule-based state, V2 representing the first wave of technological transformation, and the Current Version representing the fully mature, AI-driven state in operation today.

For each version, the report documents how the process worked, who the actors were, what problems were encountered, and what specific changes were made to address those problems. This approach reflects the iterative nature of Business Process Management (BPM), in which process improvement is understood not as a single event but as a continuous cycle of analysis, redesign, and re-evaluation. The versioned structure allows the reader to understand not only where each process ended up, but how and why it got there.

2. Process 1: KYC and Customer Onboarding

2.1 Process Overview

Know Your Customer (KYC) is a mandatory regulatory compliance process that requires financial institutions to verify the identity of their clients, understand the nature of their activities, and assess the potential risk they pose. For JPMorgan Chase, customer onboarding is the end-to-end process by which a new customer whether a retail individual, a small business, or a large multinational corporation is formally admitted into the bank's ecosystem. The process encompasses identity verification, document collection, risk scoring, compliance screening against sanctions lists, account setup, and the initiation of the banking relationship.

The process type is Application-to-Approval: a customer applies to open an account or begin a banking relationship, and the bank approves or rejects the application after completing its checks. The primary actors across all versions of this process are the Customer, the Relationship Manager, the KYC/Compliance Analyst, the Legal Team (for complex corporate cases), and in later versions automated AI systems that take over much of the analytical work previously performed by humans.

V1 Original Manual Process *(Pre-2010)*

2.2 V1: How the Process Worked

In V1, the entire KYC and customer onboarding process was manual and paper-based. When a customer applied to open an account, a Relationship Manager would meet them in person and collect physical documents passport copies, utility bills, bank statements, and for corporate clients, company registration certificates and shareholder agreements. These paper documents were then passed to a KYC analyst, who would manually read each document and physically type the extracted information into multiple disconnected internal systems: the core banking platform, the AML database, and the sanctions screening tool each operating independently with no integration between them.

Sanctions screening was performed by manually searching lists such as the OFAC Specially Designated Nationals (SDN) list and Politically Exposed Persons (PEP) databases one by one. For corporate clients, a legal team member would manually trace the ownership structure through multiple layers of subsidiaries across different jurisdictions to identify the Ultimate Beneficial Owner (UBO). Once all checks were completed and signed off by a compliance manager, the account would be activated and the customer notified by letter or phone call.

V1 Key Attributes

Attribute	Detail
Period	Pre-2010
Process type	Entirely manual and paper-based
Onboarding time retail	5 to 10 business days
Onboarding time corporate	30 to 120+ days
Sanctions screening	Manual search of individual watchlists
Data entry	Manual, repeated across disconnected systems
Ongoing monitoring	None periodic manual review only
Estimated annual cost	Hundreds of millions of USD

V1 Problems Identified

- Onboarding was extremely slow corporate clients waited months before accounts became active, causing significant customer dissatisfaction and competitive disadvantage against faster-moving competitors.
- Extremely high operational cost thousands of analysts globally performing manual, repetitive document review made this one of the most expensive compliance processes in the bank.
- Regulatory failure in 2014, JPMorgan Chase paid $2.6 billion in penalties related to its failure to detect Bernie Madoff's Ponzi scheme, with regulators directly citing failures in the manual, disconnected KYC process.
- Poor data quality manual data entry across disconnected systems produced frequent errors and inconsistencies, making a single accurate customer view impossible.
- No scalability adding more customers required hiring proportionally more analysts; the process could not keep pace with the bank's growth through acquisitions and market expansion.
- No real-time monitoring if a customer was added to a sanctions list after their initial onboarding, the bank would not detect this until the next scheduled periodic review, which could be years later.

V2 First Wave of Digitization *(2010–2018)*

2.3 V2: What Changed and How the Process Evolved

Following the regulatory failures and billion-dollar fines of the early 2010s, JPMorgan Chase made substantial investments in the first wave of digital transformation of its KYC process. The most significant structural change in V2 was the introduction of a centralized KYC utility model. Previously, each business line retail banking, investment banking, asset management had conducted its own separate KYC check on the same customer independently. In V2, a single centralized KYC operations team was established to conduct one comprehensive check, producing a shared customer profile available across all business lines. This eliminated enormous duplication of effort.

Digital document collection portals were introduced, allowing customers to upload documents electronically through secure online portals rather than submitting physical copies. Integration with commercial identity verification databases such as credit bureau APIs and national company registry services allowed certain document details to be auto-verified without human review. Automated sanctions screening tools replaced manual list searches, checking customer names against hundreds of global watchlists simultaneously within seconds and generating alerts for potential matches.

However, V2 was still substantially human-driven. While the mechanics of document collection and list screening were automated, the analytical judgment risk assessment, UBO determination for complex entities, and decision-making on borderline cases remained with human KYC analysts. Some systems still required manual data entry, and ongoing monitoring was upgraded only to an annual scheduled refresh rather than true continuous monitoring.

V2 Changes Made vs V1

Attribute	**V1**	**V2**
Document collection	Physical paper in branch	Digital upload portal
KYC check duplication	Repeated per business line	Centralized single-check model
Sanctions screening	Manual list-by-list search	Automated multi-list simultaneous screening
Risk assessment	Fully manual analyst judgment	Tool-assisted with partial scoring
Onboarding time retail	5–10 days	2–5 days

Onboarding time corporate	30–120+ days	14–45 days
Ongoing monitoring	None / irregular	Annual scheduled refresh
Data accuracy	Low (manual entry errors)	Moderate (reduced manual entry)

V2 Remaining Problems

- Risk scoring was still inconsistent different analysts might assign different risk levels to identical customers, creating compliance gaps.
- Large corporate document sets still required human review the centralized team became a new bottleneck as the customer base grew.
- Ongoing monitoring was still periodic, not continuous a customer flagged on a new sanctions list between annual reviews would not be detected.
- Cross-border compliance remained complex different regulations across 100+ countries still required significant manual handling and local expertise.
- The centralized model improved consistency but created a single point of overload the central KYC team was frequently overwhelmed during periods of high new-customer volume.

Current Version: AI-Driven Intelligent KYC *(2018–Present)*

2.4 Current Version: The Fully Evolved Process

The current version of JPMorgan Chase's KYC and onboarding process represents a fundamental redesign enabled by artificial intelligence, machine learning, and real-time data integration. The most transformative change is the deployment of natural language processing (NLP) algorithms that automatically read, classify, and extract data from customer documents including complex multi-hundred-page corporate filings in minutes rather than days. AI models trained on millions of historical documents can identify beneficial owners, financial figures, and risk indicators with accuracy comparable to experienced senior analysts.

Machine learning models now perform fully automated risk scoring for every new customer. Each customer receives a dynamic risk score based on geography, industry, transaction history, ownership structure, and adverse media signals. This score automatically determines the due diligence pathway low-risk customers are onboarded with minimal human intervention, while high-risk customers are automatically escalated to Enhanced Due Diligence (EDD) by senior compliance officers who focus exclusively on these complex cases.

Real-time biometric identity verification using facial recognition and document liveness detection allows retail customers to complete identity verification through the mobile banking application in under five minutes, with no branch visit required. The most significant conceptual shift in the current version is the move from periodic monitoring to fully continuous, real-time monitoring. Every customer's profile is checked against watchlists, adverse media sources, and corporate registry updates in real time any change triggers an immediate automated alert and review, eliminating the dangerous detection gap that existed in V1 and V2. JPMorgan Chase is also piloting blockchain-based KYC utilities that would allow verified identity data to be shared across financial institutions with customer consent.

Full Evolution Comparison KYC Process

Metric	V1 Original	V2 Transitional	Current Version
Retail onboarding time	5–10 days	2–5 days	Minutes to 1 day
Corporate onboarding time	30–120+ days	14–45 days	3–14 days
Document processing	Manual human review	Partially automated	Full AI/NLP automation
Risk scoring	Manual analyst judgment	Tool-assisted manual	Automated ML scoring
Sanctions screening	Manual list search	Automated multi-list	Real-time continuous AI
Monitoring	None / irregular	Annual scheduled	Real-time continuous
Identity verification	In-branch paper check	Digital document upload	Mobile biometric AI
Data accuracy	Low	Moderate	High
False positive rate	High	Moderate	Reduced by ~70%
Regulatory compliance posture	Reactive (post-failure)	Partially proactive	Fully proactive

3. Process 2: Fraud Detection

3.1 Process Overview

Fraud detection is among the most operationally critical processes at JPMorgan Chase. With hundreds of millions of transactions processed every day across credit cards, debit cards, wire transfers, ACH payments, and digital banking channels, the bank must identify fraudulent transactions ideally before they are authorized while ensuring that legitimate customer transactions are not incorrectly blocked. The process type is Issue-to-Resolution: a potentially fraudulent transaction is the issue, and the resolution is either confirming it as genuine and allowing it to proceed, or confirming it as fraud, blocking it, reversing any impact, and protecting the customer.

Actors across all versions include: the Customer, the Fraud Analyst, the IT/Rules Development Team, the Risk Management Team, Customer Service, and in later versions AI and ML systems operating as automated actors within the process flow.

V1 Rule-Based Static Detection *(Pre-2012)*

3.2 V1: How the Process Worked

In V1, fraud detection at JPMorgan Chase was built entirely on static rule-based systems. Fraud analysts and IT developers collaborated to write explicit if-then rules based on known fraud patterns observed in historical cases. Example rules included: 'If a transaction occurs in a foreign country within two hours of a domestic transaction, flag as suspicious,' or 'If more than five card transactions above $500 occur within one hour, block the card.' When a customer initiated a transaction, the rule engine evaluated it against all active rules. If any rule fired, the transaction was either automatically declined or queued for manual analyst review.

Fraud analysts would then manually review each flagged transaction, contacting the customer by phone if necessary to confirm whether the transaction was genuine. When new fraud patterns emerged that were not covered by existing rules, analysts would document the pattern, write a new rule specification, submit it to the IT team for development and testing, and wait for deployment a cycle that often took days or weeks. The entire process was reactive: fraud losses had to be sustained before a new pattern could be identified, encoded as a rule, and deployed.

V1 Key Attributes

Attribute	**Detail**
Period	Pre-2012
Detection method	Static if-then rules coded by analysts and IT
Variables considered per transaction	Tens (explicit rule conditions)
Detection timing	Post-transaction after money had already moved
False positive rate	Up to 50 false alerts per genuine fraud case
Novel fraud detection ability	None requires manual rule creation cycle first
Alert review method	100% manual analyst review of each flagged transaction
New rule deployment time	Days to weeks per new fraud type

V1 Problems Identified

- Extremely high false positive rate up to 50 legitimate transactions were blocked for every single genuine fraud case, causing massive customer dissatisfaction, embarrassment for customers whose transactions were declined in public, and high customer service call volumes.
- Completely reactive process new fraud techniques could not be detected until a human analyst had observed them, written a rule, had it tested by IT, and deployed it. Fraudsters had a guaranteed window of opportunity every time they changed techniques.
- Alert fatigue among analysts the overwhelming volume of false positive alerts meant genuine fraud was sometimes missed because analysts were buried in noise.
- Inability to detect novel fraud synthetic identity fraud, which began emerging in the mid-2000s by combining real and fabricated personal information to create fictional identities, was completely undetectable because it matched no existing rule pattern.
- Scalability failure as transaction volumes grew exponentially with online and mobile banking adoption, the rule engine struggled to evaluate thousands of rules against millions of transactions per minute without introducing latency into transaction authorization.
- Significant annual fraud losses sustained because the system could only ever detect patterns that had already been seen, understood, and manually encoded.

V2 First Generation Machine Learning *(2012–2019)*

3.3 V2: What Changed and How the Process Evolved

Following the 2012 London Whale trading scandal which, while not traditional fraud, exposed deep weaknesses in JPMorgan Chase's risk monitoring capabilities and cost the bank over $6 billion the bank made major investments in data analytics and machine learning for risk and fraud detection. V2 introduced a hybrid architecture in which machine learning models ran alongside the existing rule-based system rather than replacing it outright.

ML models were trained on massive historical transaction datasets containing billions of past transactions both fraudulent and legitimate and learned to distinguish fraud from genuine activity based on hundreds of variables simultaneously: transaction amount, time of day, merchant category, geographic location, device type, the customer's historical spending patterns, and the velocity of recent transactions. Each transaction was assigned a fraud probability score in near-real time (within a few seconds). Transactions above a defined threshold were automatically declined; those in a middle band were queued for analyst review; low-scoring transactions proceeded without interruption.

Critically, ML models could be periodically retrained on new data typically weekly or monthly allowing them to adapt to new fraud patterns far more quickly than the old rule-writing cycle. The rule-based layer was retained in parallel for high-confidence categorical rules such as blocking transactions from sanctioned jurisdictions.

V2 Changes Made vs V1

Attribute	V1	V2
Primary detection method	Static rules only	Hybrid: ML models + retained core rules
Variables per transaction	Tens (rule conditions)	Hundreds (ML feature set)
Detection timing	Post-transaction	Near-real-time (seconds)
False positive rate	~50x per genuine case	Reduced by approximately 30–40%
Adaptation to new fraud	Weeks (manual rule cycle)	Days (model retraining)
Analyst workload	100% manual review	ML pre-filters analysts review escalations
Novel fraud detection	Not possible	Improved patterns learned from data

Account takeover detection	Very poor	Moderate improvement

V2 Remaining Problems

- Models were retrained periodically (weekly or monthly) rather than continuously a meaningful gap still existed between a new fraud technique first appearing and the model updating to detect it.
- Processing speed was near-real-time but not yet true millisecond-speed real-time, meaning the scoring system added minor latency to transaction authorization flows.
- Behavioral biometrics using the way a customer physically interacts with their device as an identity verification signal were not yet integrated, leaving account takeover fraud difficult to detect even with ML scoring.
- The parallel rule-based layer created system complexity and occasional conflicts where a rule-based block and an ML-based approval contradicted each other.
- Model explainability was limited when a transaction was declined by the ML system, providing a clear explanation to the customer or regulator was difficult.

Current Version: Real-Time AI and Behavioral Analytics *(2019–Present)*

3.4 Current Version: The Fully Evolved Process

The current version of JPMorgan Chase's fraud detection process evaluates every single transaction in under 100 milliseconds fast enough to operate invisibly within the normal card authorization flow, with zero perceptible added latency for the customer. The static rule-based layer has been almost entirely retired and replaced by advanced ensemble ML models that are continuously retrained on streaming real-time transaction data. This continuous learning means the system can adapt to a new fraud technique within hours of its first appearance rather than the days or weeks required in V2.

Behavioral analytics and biometrics now form a critical and distinct layer of the detection architecture. The bank builds a detailed behavioral profile for every customer over time, capturing patterns such as typical transaction amounts, geographic patterns, spending categories, preferred merchants, and even the precise way the customer holds their phone, types on the keyboard, and scrolls through the banking application. When a transaction or login attempt deviates significantly from this behavioral fingerprint even if the correct password has been entered the system triggers enhanced scrutiny, enabling reliable detection of account takeover fraud that was virtually invisible to V1 and V2.

Graph analytics technology maps the relationships between customers, accounts, transactions, devices, and IP addresses across the entire customer base, identifying organized fraud ring patterns that are completely invisible at the individual transaction level. A network of accounts that appear unrelated when examined individually may reveal a clear fraud ring structure when examined as a graph. JPMorgan Chase also participates in industry-wide fraud intelligence sharing networks, where anonymized fraud signals detected at one institution are rapidly shared across the network, allowing all participating banks to update their defenses within hours of a new fraud technique being first observed anywhere in the financial system.

Human fraud analysts continue to be essential but their role has been fundamentally redefined. Rather than spending their day reviewing hundreds of routine flagged transactions as in V1 they now focus exclusively on complex, high-value, ambiguous cases that genuinely require human judgment. AI tools automatically compile all relevant context for each case full account history, device data, geographic patterns, and similar historical cases and present it to the analyst in a consolidated dashboard, enabling investigative decisions in minutes that previously required hours of manual research.

Full Evolution Comparison Fraud Detection Process

Metric	V1 Rule-Based	V2 First ML	Current Version
Detection method	Static if-then rules	Hybrid ML + core rules	Full real-time AI + behavioral
Decision speed	Seconds (rule evaluation)	Near-real-time (seconds)	< 100 milliseconds
False positive rate	~50x per genuine fraud	Reduced ~30–40%	Reduced ~60–70%
Novel fraud detection	Not possible	Days after emergence	Hours after emergence
Account takeover detection	Very poor	Moderate	Strong (behavioral biometrics)
Fraud ring detection	Not possible	Limited	Full graph analytics
Model adaptation speed	Weeks (manual rule writing)	Days (periodic retraining)	Hours (continuous learning)

Human analyst role	Reviews every flagged transaction	Reviews ML-escalated cases	Investigates complex cases only
Customer experience	Frequent false declines	Fewer false declines	Minimal false declines
Annual fraud loss trend	High baseline losses	Moderate reduction	20–40% overall reduction

4. Comparative Analysis of Both Processes

Both the KYC/Customer Onboarding process and the Fraud Detection process followed a remarkably similar evolutionary trajectory across their three versions. In both cases, V1 represented a manual or rule-based state designed for an earlier, lower-volume era that became increasingly overwhelmed by scale, complexity, and external pressures. V2 represented a pragmatic first-generation technological response that delivered meaningful improvements without fully resolving the underlying structural limitations. The current version represents a genuine paradigm shift, in which artificial intelligence and real-time data processing enable capabilities that were simply not conceivable in the V1 era.

Dimension	**KYC / Onboarding**	**Fraud Detection**
V1 core problem	Speed, cost, regulatory failure from manual processes	High false positives, reactive detection, no novel pattern detection
V1 to V2 key change	Centralized model + digital portals + automated list screening	ML models introduced alongside rule engine in hybrid architecture
V2 to Current key change	Full AI document processing + continuous monitoring + biometric ID	Real-time behavioral AI + graph analytics + continuous model learning
Primary driver of evolution	Regulatory fines ($2.6B) and competitive pressure	Rising fraud losses and exponential transaction volume growth
Nature of human role today	EDD specialists review escalated high-risk cases only	Complex fraud investigators handle ambiguous high-value cases only
Remaining challenge	Cross-border data privacy regulation	Adversarial AI-powered fraud evolution

A common thread across both evolutions is the progressive compression of the human role in routine processing and the elevation of human expertise to higher-value judgment work. In V1, human analysts performed almost all of the work in both processes. In V2, automation handled the most mechanical tasks while humans retained analytical responsibility. In the current version, AI handles virtually all routine analytical work, with humans reserved for cases that genuinely require contextual judgment, ethical reasoning, or regulatory accountability.

5. Conclusion

The versioned analysis of JPMorgan Chase's KYC/Customer Onboarding and Fraud Detection processes demonstrates that successful business process evolution is not a single transformation event but a deliberate, iterative journey driven by a clear understanding of root causes. In both cases, V1 was a baseline process designed for a simpler era overwhelmed by growth, regulatory complexity, and the emergence of new threats. V2 represented a valuable but incomplete technological response, introducing automation at the mechanical level while leaving the analytical core largely unchanged. The current version represents a true paradigm shift, with AI and real-time data analytics delivering capabilities that are qualitatively not just quantitatively different from what came before.

JPMorgan Chase did not simply hire more analysts when onboarding slowed down it identified that the root cause was the manual, paper-based, disconnected structure of the process itself and redesigned it from the ground up. It did not simply write more fraud rules when detection rates fell it recognized that the rule-based paradigm was structurally incapable of keeping pace with adaptive adversaries and replaced it with a fundamentally different approach built on continuous learning.

For students of Business Process Engineering, the versioned evolution of JPMorgan Chase's processes illustrates that the most impactful process transformations are those that challenge and redesign the underlying logic of the process not merely the tools used to execute it. V1 to V2 improved the tools. V2 to Current changed the logic. Understanding this distinction is central to effective business process management.

6. References

1. JPMorgan Chase & Co. (2023). Annual Report 2023. JPMorgan Chase & Co.

2. Financial Action Task Force. (2023). Guidance on Digital Identity. FATF/OECD.

3. Thomson Reuters. (2017). Know Your Customer Surveys: Costs, Complexity and Future Outlook.

4. U.S. Department of Justice. (2014). JPMorgan Chase Agrees to Pay $2.6 Billion in Connection with Madoff Fraud.

5. U.S. Department of Justice. (2013). Justice Department Secures Record $13 Billion Settlement with JPMorgan.

6. Dimon, J. (2022). Letter to Shareholders. JPMorgan Chase & Co. Annual Report 2022.

7. McKinsey & Company. (2022). The Future of KYC: Transforming Compliance Operations. McKinsey Global Institute.

8. Accenture. (2021). Transforming Financial Crime Compliance: The AI-Powered Bank.

9. PricewaterhouseCoopers. (2023). Global Financial Crime Survey. PwC Financial Services.

10. Dumas, M., La Rosa, M., Mendling, J., & Reijers, H. A. (2018). Fundamentals of Business Process Management (2nd ed.). Springer.

11. Basel Committee on Banking Supervision. (2021). Sound Management of Risks Related to Money Laundering. Bank for International Settlements.

12. Javelin Strategy & Research. (2023). Identity Fraud Study: The Fraud Shift. Javelin Strategy & Research.